

Eric Anthony Mitchell

eric.mitchell@cs.stanford.edu / ericmitchell.ai

Education

Stanford University, Stanford, CA 2019 - 2024 (Expected)
Ph.D. Candidate, Computer Science
Advisors: Chelsea Finn & Christopher D. Manning

Princeton University, Princeton, NJ 2014 - 2018
B.S.E., Computer Science (Highest Honors)
Advisor: H. Sebastian Seung

Professional Experience

DeepMind (AlphaCode team) *London, England*
Research Scientist Intern [hosts: Junyoung Chung, Nate Kushman, Aäron van den Oord] 2022

Samsung Research America *New York City, New York*
Research Intern, part-time 2019 - 2020
Machine Learning Research Engineer 2018 - 2019

Princeton Neuroscience Institute *Princeton, New Jersey*
Research Engineer 2018

Google *Mountain View, California*
Software Engineering Intern 2017

TapTools LLC *Winston-Salem, North Carolina*
Founder, iOS Developer 2012 - 2016

Selected Works

**indicates equal contribution*

Katherine Tian*, **Eric Mitchell***, Huaxiu Yao, Christopher D. Manning, Chelsea Finn. 2023
Fine-tuning Language Models for Factuality. *Preprint*.

Eric Mitchell, Rafael Rafailov, Archit Sharma, Chelsea Finn, Christopher D. Manning. An 2023
Emulator for Fine-Tuning Large Language Models using Small Language Models. *Preprint*.

Rafael Rafailov*, Archit Sharma*, **Eric Mitchell***, Stefano Ermon, Christopher D. Manning, 2023
Chelsea Finn. Direct Preference Optimization: Your Language Model is Secretly a Reward Model.
Oral Presentation (<1% of submissions), *Neural Information Processing Systems (NeurIPS)*.

Nathan Hu*, **Eric Mitchell***, Christopher D. Manning, Chelsea Finn. Meta-Learning Online 2023
Adaptation of Language Models. *Empirical Methods in Natural Language Processing (EMNLP)*.

Katherine Tian*, **Eric Mitchell***, Allan Zhou, Archit Sharma, Rafael Rafailov, Huaxiu Yao, 2023
Chelsea Finn, Christopher D. Manning. Just Ask for Calibration: Strategies for Eliciting
Calibrated Confidence Scores from Language Models Fine-Tuned with Human Feedback.
Empirical Methods in Natural Language Processing (EMNLP).

Eric Mitchell, Yoonho Lee, Alexander Khazatsky, Christopher D. Manning, Chelsea Finn. 2023
DetectGPT: DetectGPT: Zero-Shot Machine-Generated Text Detection using Probability
Curvature. **Oral Presentation (2% of submissions)**, *International Conference on Machine
Learning (ICML)*.

Peter Henderson*, **Eric Mitchell***, Christopher D. Manning, Dan Jurafsky, Chelsea Finn. 2023
Self-Destructing Models: Increasing the Costs of Harmful Dual Uses in Foundation Models.
Honorable Mention, Best Student Paper. *AAAI/ACM Conference on Artificial Intelligence,
Ethics, & Society*.

Eric Mitchell , Joseph J Noh, Siyan Li, William S Armstrong, Ananth Agarwal, Patrick Liu, Chelsea Finn, Christopher D Manning. Enhancing Self-Consistency and Performance of Pretrained Language Models with NLI. Oral Presentation (4% of submissions) , <i>Empirical Methods in Natural Language Processing (EMNLP)</i> .	2022
Eric Mitchell , Charles Lin, Antoine Bosselut, Christopher D. Manning, Chelsea Finn. Memory-Based Model Editing at Scale. <i>International Conference on Machine Learning (ICML)</i> .	2022
Eric Mitchell , Charles Lin, Antoine Bosselut, Chelsea Finn, Christopher D. Manning. Fast Model Editing at Scale. <i>International Conference on Learning Representations (ICLR)</i> .	2022
Eric Mitchell , Rafael Rafailov, Xue Bin Peng, Sergey Levine, and Chelsea Finn. Offline Meta-Reinforcement Learning with Advantage Weighting. <i>International Conference on Machine Learning (ICML)</i> .	2021
Eric Mitchell , Selim Engin, Volkan Isler, Daniel D Lee. Higher-Order Function Networks for Learning Composable 3D Object Representations. <i>International Conference on Learning Representations (ICLR)</i> .	2020
Tarik Tosun*, Eric Mitchell* , Ben Eisner, Jinwook Huh, Bhoram Lee, Daewon Lee, Volkan Isler, H. Sebastian Seung, Daniel D Lee. Pixels to Plans: Learning Non-Prehensile Manipulation by Imitating a Planner. <i>IEEE International Conference on Intelligent Robots and Systems (IROS)</i> .	2019

Academic Awards

Stanford Knight-Hennessy Graduate Fellowship	2019 - 2022
Princeton CS Outstanding Undergraduate Thesis Prize	2018
Princeton CS Best Undergraduate Research Poster Prize	2017
Srixon/Cleveland Golf All-America Scholar	2017
Tau Beta Pi Engineering Honor Society Vice President	2016 - 2017
Princeton Shapiro Prize for Academic Excellence	2016

Service

NeurIPS Socially Responsible Language Modeling Workshop Reviewer	2023
NeurIPS Conference Reviewer	2020 - 2021, 2023
Stanford CS Ph.D. Admissions Committee Student Representative	2020 - 2022
ICLR Conference Reviewer	2020 - 2021
NeurIPS Meta-Learning Workshop Reviewer	2020 - 2021
AI4ALL Summer Camp Instructor	2018

Other Activities

PADI Open Water SCUBA certification	2023
Princeton NCAA Division I Varsity Golf Team	2014 - 2018
HackPrinceton (1 st Place Team, Software)	2015
Guitar, Singing, Songwriting	